

GAURAV SAHU

GITHUB LINK: <https://github.com/GauravSahu13>

IBM HR EMPLOYEE ATTRITION

```
In [2]: import pandas as pd
import seaborn as sns
import numpy as np
from statistics import mean
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")
%matplotlib inline
```

```
In [22]: !pip install wget

Collecting wget
  Downloading wget-3.2.zip (10 kB)
  Preparing metadata (setup.py) ... done
Building wheels for collected packages: wget
  Building wheel for wget (setup.py) ... done
  Created wheel for wget: filename=wget-3.2-py3-none-any.whl size=9655 sha256=85a0ea154a55e0b27af40c2394a7f138f6a082774359f02a038b7806dc11f98f
  Stored in directory: /root/.cache/pip/wheels/8b/f1/7f/5c94f0a7a505ca1c81cd1d9208ae2064675d97582078e6c769
Successfully built wget
Installing collected packages: wget
Successfully installed wget-3.2
```

PreviewCodeBlame

RawCopyDownloadEditDropdown

```
# Load csv file
df = pd.read_csv('HR-Employee-Attrition.csv')
df.head()
```

Out[23]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Education
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sci
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sci
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sci
4	27	No	Travel_Rarely	591	Research & Development	2	1	M

5 rows × 35 columns

```
In [24]: df.shape
```

Out[24]: (1470, 35)

```
In [25]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                           1470 non-null   object
2   BusinessTravel                      1470 non-null   object
3   DailyRate                           1470 non-null   int64
4   Department                          1470 non-null   object
5   DistanceFromHome                   1470 non-null   int64
6   Education                           1470 non-null   int64
7   EducationField                      1470 non-null   object
8   EmployeeCount                      1470 non-null   int64
9   EmployeeNumber                     1470 non-null   int64
10  EnvironmentSatisfaction             1470 non-null   int64
11  Gender                             1470 non-null   object
12  HourlyRate                         1470 non-null   int64
13  JobInvolvement                     1470 non-null   int64
14  JobLevel                           1470 non-null   int64
15  JobRole                            1470 non-null   object
16  JobSatisfaction                     1470 non-null   int64
17  MaritalStatus                      1470 non-null   object
18  MonthlyIncome                      1470 non-null   int64
19  MonthlyRate                        1470 non-null   int64
20  NumCompaniesWorked                 1470 non-null   int64
21  Over18                             1470 non-null   object
22  OverTime                           1470 non-null   object
23  PercentSalaryHike                  1470 non-null   int64
24  PerformanceRating                  1470 non-null   int64
25  RelationshipSatisfaction            1470 non-null   int64
26  StandardHours                      1470 non-null   int64
27  StockOptionLevel                   1470 non-null   int64
28  TotalWorkingYears                  1470 non-null   int64
29  TrainingTimesLastYear              1470 non-null   int64
30  WorkLifeBalance                    1470 non-null   int64
31  YearsAtCompany                     1470 non-null   int64
32  YearsInCurrentRole                 1470 non-null   int64
33  YearsSinceLastPromotion             1470 non-null   int64
34  YearsWithCurrManager                1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

```
In [26]: df.isnull().sum()
```

```
Out[26]: Age                                0
Attrition                                0
BusinessTravel                          0
DailyRate                              0
Department                              0
DistanceFromHome                        0
Education                               0
EducationField                          0
EmployeeCount                           0
EmployeeNumber                          0
EnvironmentSatisfaction                  0
Gender                                  0
HourlyRate                              0
JobInvolvement                          0
JobLevel                                0
JobRole                                 0
JobSatisfaction                         0
MaritalStatus                           0
MonthlyIncome                           0
MonthlyRate                             0
NumCompaniesWorked                      0
Over18                                  0
OverTime                                0
PercentSalaryHike                       0
PerformanceRating                       0
RelationshipSatisfaction                  0
StandardHours                           0
```

```

StockOptionLevel      0
TotalWorkingYears     0
TrainingTimesLastYear  0
WorkLifeBalance        0
YearsAtCompany         0
YearsInCurrentRole     0
YearsSinceLastPromotion 0
YearsWithCurrManager   0
dtype: int64

```

In [27]: `df.nunique()`

```

Out[27]: Age          43
Attrition           2
BusinessTravel      3
DailyRate          886
Department          3
DistanceFromHome    29
Education            5
EducationField       6
EmployeeCount        1
EmployeeNumber     1470
EnvironmentSatisfaction 4
Gender              2
HourlyRate          71
JobInvolvement       4
JobLevel             5
JobRole              9
JobSatisfaction      4
MaritalStatus        3
MonthlyIncome       1349
MonthlyRate         1427
NumCompaniesWorked   10
Over18               1
OverTime             2
PercentSalaryHike    15
PerformanceRating    2
RelationshipSatisfaction 4
StandardHours        1
StockOptionLevel     4
TotalWorkingYears    40
TrainingTimesLastYear 7
WorkLifeBalance      4
YearsAtCompany       37
YearsInCurrentRole   19
YearsSinceLastPromotion 16
YearsWithCurrManager 18
dtype: int64

```

In [28]: `df.duplicated().sum()`

Out[28]: 0

In [29]: `df.describe().T`

```

Out[29]:
```

	count	mean	std	min	25%	50%	75%	max
Age	1470.0	36.923810	9.135373	18.0	30.00	36.0	43.00	60
DailyRate	1470.0	802.485714	403.509100	102.0	465.00	802.0	1157.00	1499
DistanceFromHome	1470.0	9.192517	8.106864	1.0	2.00	7.0	14.00	29
Education	1470.0	2.912925	1.024165	1.0	2.00	3.0	4.00	5
EmployeeCount	1470.0	1.000000	0.000000	1.0	1.00	1.0	1.00	1
EmployeeNumber	1470.0	1024.865306	602.024335	1.0	491.25	1020.5	1555.75	2068
EnvironmentSatisfaction	1470.0	2.721769	1.093082	1.0	2.00	3.0	4.00	4
HourlyRate	1470.0	65.891156	20.329428	30.0	48.00	66.0	83.75	100

JobInvolvement	1470.0	2.729932	0.711561	1.0	2.00	3.0	3.00	4
JobLevel	1470.0	2.063946	1.106940	1.0	1.00	2.0	3.00	5
JobSatisfaction	1470.0	2.728571	1.102846	1.0	2.00	3.0	4.00	4
MonthlyIncome	1470.0	6502.931293	4707.956783	1009.0	2911.00	4919.0	8379.00	19999
MonthlyRate	1470.0	14313.103401	7117.786044	2094.0	8047.00	14235.5	20461.50	26999
NumCompaniesWorked	1470.0	2.693197	2.498009	0.0	1.00	2.0	4.00	9
PercentSalaryHike	1470.0	15.209524	3.659938	11.0	12.00	14.0	18.00	25
PerformanceRating	1470.0	3.153741	0.360824	3.0	3.00	3.0	3.00	4
RelationshipSatisfaction	1470.0	2.712245	1.081209	1.0	2.00	3.0	4.00	4
StandardHours	1470.0	80.000000	0.000000	80.0	80.00	80.0	80.00	80
StockOptionLevel	1470.0	0.793878	0.852077	0.0	0.00	1.0	1.00	3
TotalWorkingYears	1470.0	11.279592	7.780782	0.0	6.00	10.0	15.00	40
TrainingTimesLastYear	1470.0	2.799320	1.289271	0.0	2.00	3.0	3.00	6
WorkLifeBalance	1470.0	2.761224	0.706476	1.0	2.00	3.0	3.00	4
YearsAtCompany	1470.0	7.008163	6.126525	0.0	3.00	5.0	9.00	40
YearsInCurrentRole	1470.0	4.229252	3.623137	0.0	2.00	3.0	7.00	18
YearsSinceLastPromotion	1470.0	2.187755	3.222430	0.0	0.00	1.0	3.00	15
YearsWithCurrManager	1470.0	4.123129	3.568136	0.0	2.00	3.0	7.00	17

In [30]: `df.columns`


Out[30]: Index(['Age', 'Attrition', 'BusinessTravel', 'DailyRate', 'Department', 'DistanceFromHome', 'Education', 'EducationField', 'EmployeeCount', 'EmployeeNumber', 'EnvironmentSatisfaction', 'Gender', 'HourlyRate', 'JobInvolvement', 'JobLevel', 'JobRole', 'JobSatisfaction', 'MaritalStatus', 'MonthlyIncome', 'MonthlyRate', 'NumCompaniesWorked', 'Over18', 'OverTime', 'PercentSalaryHike', 'PerformanceRating', 'RelationshipSatisfaction', 'StandardHours', 'StockOptionLevel', 'TotalWorkingYears', 'TrainingTimesLastYear', 'WorkLifeBalance', 'YearsAtCompany', 'YearsInCurrentRole', 'YearsSinceLastPromotion', 'YearsWithCurrManager'], dtype='object')


In all we have 34 features consisting of both the categorical as well as the numerical features. The target variable is the 'Attrition' of the employee which can be either a Yes or a No.

Univariate Analysis

In [31]: `sns.kdeplot(df['Age'], shade=True, color='#ff4125')`


Out[31]: <Axes: xlabel='Age', ylabel='Density'>







In [32]: `sns.distplot(df['Age'])`

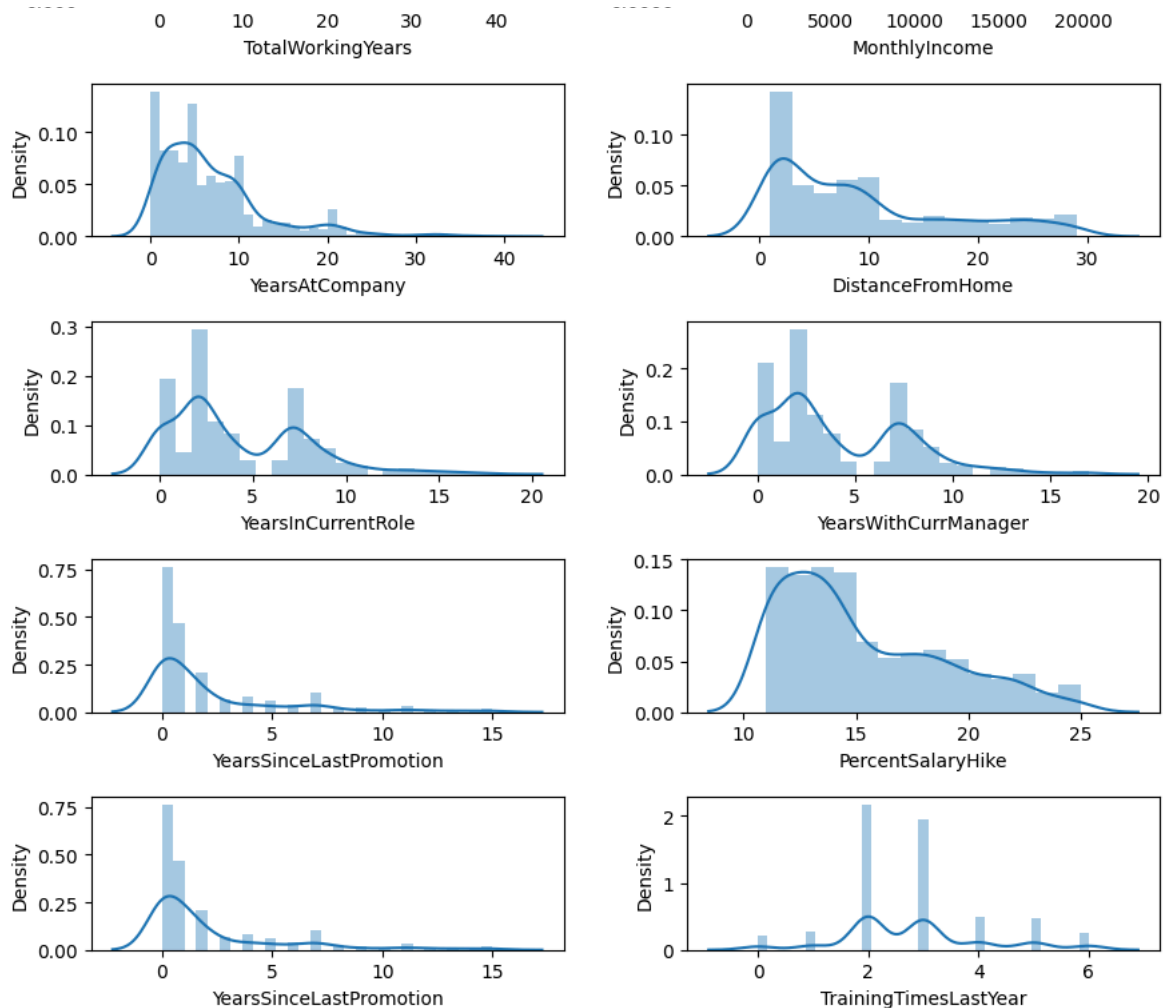
Out[32]: `<Axes: xlabel='Age', ylabel='Density'>`



In [33]: `warnings.filterwarnings('always')`
`warnings.filterwarnings('ignore')`

```
fig,ax = plt.subplots(5,2, figsize=(9,9))
sns.distplot(df['TotalWorkingYears'], ax = ax[0,0])
sns.distplot(df['MonthlyIncome'], ax = ax[0,1])
sns.distplot(df['YearsAtCompany'], ax = ax[1,0])
sns.distplot(df['DistanceFromHome'], ax = ax[1,1])
sns.distplot(df['YearsInCurrentRole'], ax = ax[2,0])
sns.distplot(df['YearsWithCurrManager'], ax = ax[2,1])
sns.distplot(df['YearsSinceLastPromotion'], ax = ax[3,0])
sns.distplot(df['PercentSalaryHike'], ax = ax[3,1])
sns.distplot(df['YearsSinceLastPromotion'], ax = ax[4,0])
sns.distplot(df['TrainingTimesLastYear'], ax = ax[4,1])
plt.tight_layout()
plt.show()
```






```
In [34]: cat_df=df.select_dtypes(include='object')
cat_df.columns
```

```
Out[34]: Index(['Attrition', 'BusinessTravel', 'Department', 'EducationField', 'Gender',
               'JobRole', 'MaritalStatus', 'Over18', 'OverTime'],
              dtype='object')
```

```
In [35]: sns.set_style('darkgrid')
sns.countplot(x='Attrition', data = df)
```

```
Out[35]: <Axes: xlabel='Attrition', ylabel='count'>
```




Yes

No

Attrition


```
In [36]: sns.lmplot(x = 'Age', y = 'DailyRate', hue = 'Attrition', data = df)
```

```
Out[36]: <seaborn.axisgrid.FacetGrid at 0x7ef6c7797400>
```




```
In [37]: plt.figure(figsize =(10, 6))
sns.boxplot(y ='MonthlyIncome', x ='Attrition', data = df)
```

```
Out[37]: <Axes: xlabel='Attrition', ylabel='MonthlyIncome'>
```



```
In [38]: sns.countplot(x ='OverTime', data = df)
```


```
Out[38]: <Axes: xlabel='OverTime', ylabel='count'>
```



```
In [39]:
```

```
sns.countplot(x='Gender', data=df)
```


```
Out[39]: <Axes: xlabel='Gender', ylabel='count'>
```




```
In [40]:
```

```
fig, ax = plt.subplots(figsize=(20, 5))  
sns.countplot(x='JobRole', data=df)
```


```
Out[40]: <Axes: xlabel='JobRole', ylabel='count'>
```





```
In [41]: #corelation matrix.
cor_mat= df.corr()
mask = np.array(cor_mat)
mask[np.tril_indices_from(mask)] = False
fig=plt.gcf()
fig.set_size_inches(30,12)
sns.heatmap(data=cor_mat,mask=mask,square=True,annot=True,cbar=True)
```

Out[41]: <Axes: >



```
In [42]: # Feature Selection
df.drop(['BusinessTravel', 'DailyRate', 'EmployeeCount', 'EmployeeNumber', 'HourlyRate', 'Monthly',
        , 'NumCompaniesWorked', 'Over18', 'StandardHours', 'StockOptionLevel', 'TrainingTimes
```

```
In [43]: df.shape
```

Out[43]: (1470, 24)

Feature Encoding

```
In [44]: #import the necessary modelling algos.
from sklearn.linear_model import LogisticRegression
from sklearn.svm import LinearSVC
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
```

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.naive_bayes import GaussianNB

#model selection
from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold
from sklearn.metrics import accuracy_score, precision_score, recall_score, confusion_matrix, roc
from sklearn.model_selection import GridSearchCV

from imblearn.over_sampling import SMOTE

#preprocess.
from sklearn.preprocessing import MinMaxScaler, StandardScaler, LabelEncoder, OneHotEncoder

```

```

In [45]: def transform(feature):
          le=LabelEncoder()
          df[feature]=le.fit_transform(df[feature])
          print(le.classes_)

```

```

In [46]: cat_df=df.select_dtypes(include='object')
          cat_df.columns

```

```

Out[46]: Index(['Attrition', 'Department', 'EducationField', 'Gender', 'JobRole',
               'MaritalStatus', 'OverTime'],
              dtype='object')

```

```

In [47]: for col in cat_df.columns:
          transform(col)

```

```

['No' 'Yes']
['Human Resources' 'Research & Development' 'Sales']
['Human Resources' 'Life Sciences' 'Marketing' 'Medical' 'Other'
 'Technical Degree']
['Female' 'Male']
['Healthcare Representative' 'Human Resources' 'Laboratory Technician'
 'Manager' 'Manufacturing Director' 'Research Director'
 'Research Scientist' 'Sales Executive' 'Sales Representative']
['Divorced' 'Married' 'Single']
['No' 'Yes']

```

```

In [48]: y = df.iloc[:, 1]
          X = df
          X.drop('Attrition', axis = 1, inplace = True)

```

```

In [49]: X.head()


```

```

Out[49]:
   Age  Department  DistanceFromHome  Education  EducationField  EnvironmentSatisfaction  Gender
0    41           2                   1          2                1                      2         0
1    49           1                   8          1                1                      3         1
2    37           1                   2          2                4                      4         1
3    33           1                   3          4                1                      4         0
4    27           1                   2          1                3                      1         1

```

5 rows × 23 columns



```

In [50]: #Splitting data into Training and Test data
          X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.15, random_state=20)

```

```
In [51]: X_train.shape, y_train.shape
```

```
Out[51]: ((1249, 23), (1249,))
```

```
In [52]: ### Crating a standard scaler object
scaler=StandardScaler()
scaler
```

```
Out[52]: StandardScaler()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [53]: ### using fit_transform to Standardize the train data
X_train=scaler.fit_transform(X_train)
X_train
```

```
Out[53]: array([[ 1.41518706, -0.5031381 ,  0.9840021 , ...,  1.835334 ,
                -0.38987004,  1.92450052],
                [-0.21943252, -0.5031381 ,  1.97915888, ...,  1.56178601,
                 2.34238001, -0.88651708],
                [-0.87328035,  1.40695153,  1.73036969, ...,  1.01469003,
                 0.21729664,  1.64339876],
                ...,
                [ 1.08826314, -0.5031381 ,  2.47673727, ...,  0.74114204,
                 2.03879667,  1.362297  ],
                [-1.41815354, -0.5031381 ,  1.73036969, ...,  0.19404605,
                 -0.38987004, -0.0432118 ],
                [-1.09122963, -2.41322773,  1.60597509, ..., -1.1736939 ,
                 -0.69345337, -1.16761884]])
```

```
In [54]: ### here using transform only to avoid data leakage
### (training mean and training std will be used for standardisation when we use transform)
X_test=scaler.transform(X_test)
X_test
```

```
Out[54]: array([[ 0.1074914 , -0.5031381 ,  1.1083967 , ..., -1.1736939 ,
                -0.38987004, -1.16761884],
                [-0.65533107, -0.5031381 , -0.01115468, ...,  1.01469003,
                 -0.69345337,  1.08119524],
                [-0.4373818 , -0.5031381 , -0.50873307, ...,  0.46759404,
                 1.73521333,  1.08119524],
                ...,
                [-0.54635643,  1.40695153, -0.50873307, ...,  1.01469003,
                 0.82446332, -0.32431356],
                [ 0.65236459, -0.5031381 ,  2.35234268, ..., -0.62659792,
                 -0.38987004, -0.60541532],
                [ 0.54338995, -0.5031381 , -0.88191686, ...,  0.46759404,
                 0.52087998,  2.7678058 ]])
```

```
In [55]: from sklearn.neighbors import KNeighborsClassifier
neighbors = []
cv_scores = []

from sklearn.model_selection import cross_val_score
# perform 10 fold cross validation
for k in range(1, 40, 2):
    neighbors.append(k)
    knn = KNeighborsClassifier(n_neighbors = k)
    scores = cross_val_score(
        knn, X_train, y_train, cv = 10, scoring = 'accuracy')
    cv_scores.append(scores.mean())
error_rate = [1-x for x in cv_scores]

# determining the best k
optimal_k = neighbors[error_rate.index(min(error_rate))]
print(f'The optimal number of neighbors is %d ! % optimal k'
```

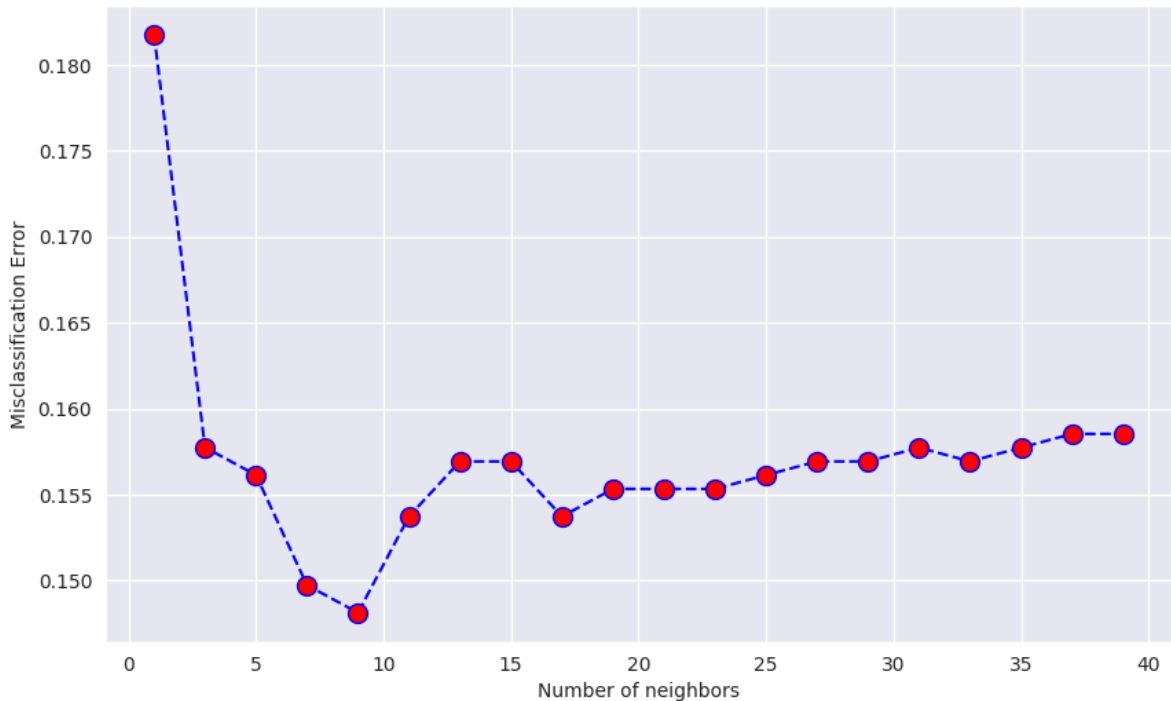
```

print( 'The optimal number of neighbors is % d' % optimal_k)

# plot misclassification error versus k
plt.figure(figsize = (10, 6))
plt.plot(range(1, 40, 2), error_rate, color = 'blue', linestyle = 'dashed', marker = 'o',
         markerfacecolor = 'red', markersize = 10)
plt.xlabel('Number of neighbors')
plt.ylabel('Misclassification Error')
plt.show()

```

The optimal number of neighbors is 9



```

In [56]: from sklearn.model_selection import cross_val_predict, cross_val_score
from sklearn.metrics import accuracy_score, classification_report
from sklearn.metrics import confusion_matrix

def print_score(clf, X_train, y_train, X_test, y_test, train = True):
    if train:
        print("Train Result:")
        print("-----")
        print("Classification Report: \n {}".format(classification_report(
            y_train, clf.predict(X_train))))
        print("Confusion Matrix: \n {}".format(confusion_matrix(
            y_train, clf.predict(X_train))))

        res = cross_val_score(clf, X_train, y_train,
                               cv = 10, scoring = 'accuracy')
        print("Average Accuracy: \t {}".format(np.mean(res)))
        print("Accuracy SD: \t {}".format(np.std(res)))
        print("accuracy score: {}".format(accuracy_score(
            y_train, clf.predict(X_train))))
        print("-----")

    elif train == False:
        print("Test Result:")
        print("-----")
        print("Classification Report: \n {}".format(
            classification_report(y_test, clf.predict(X_test))))
        print("Confusion Matrix: \n {}".format(
            confusion_matrix(y_test, clf.predict(X_test))))
        print("accuracy score: {}".format(
            accuracy_score(y_test, clf.predict(X_test))))
        print("-----")

knn = KNeighborsClassifier(n_neighbors = 7)

```

```
knn.fit(X_train, y_train)
print_score(knn, X_train, y_train, X_test, y_test, train = True)
print_score(knn, X_train, y_train, X_test, y_test, train = False)
```

Train Result:

Classification Report:

	precision	recall	f1-score	support
0	0.87	1.00	0.93	1047
1	0.90	0.22	0.35	202
accuracy			0.87	1249
macro avg	0.88	0.61	0.64	1249
weighted avg	0.87	0.87	0.83	1249

Confusion Matrix:

```
[[1042  5]
 [ 158 44]]
```

Average Accuracy: 0.8503
 Accuracy SD: 0.0103
 accuracy score: 0.8695

Test Result:

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.99	0.92	186
1	0.60	0.09	0.15	35
accuracy			0.85	221
macro avg	0.73	0.54	0.53	221
weighted avg	0.81	0.85	0.79	221

Confusion Matrix:

```
[[184  2]
 [ 32  3]]
```

accuracy score: 0.8462

In [57]:

```
# Validation scores of all base models
from sklearn.preprocessing import scale, StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.metrics import confusion_matrix, accuracy_score, mean_squared_error, r2_score,
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from lightgbm import LGBMClassifier
from sklearn.model_selection import KFold

models = []
models.append(('Log', LogisticRegression()))
models.append(('KNN', KNeighborsClassifier()))
models.append(('CART', DecisionTreeClassifier(random_state = 12345)))
models.append(('RF', RandomForestClassifier(random_state = 12345)))
models.append(('SVM', SVC(gamma='auto', random_state = 12345)))
models.append(('XGB', GradientBoostingClassifier(random_state = 12345)))
models.append(('LightGBM', LGBMClassifier(random_state = 12345)))

# evaluate each model in turn
results = []
names = []
```

```
In [58]: for name, model in models:

        kfold = KFold(n_splits = 10, random_state = 12345, shuffle=True)
        cv_results = cross_val_score(model, X, y, cv = 10, scoring= "accuracy")
        results.append(cv_results)
        names.append(name)
        msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
        print(msg)

# boxplot algorithm comparison
fig = plt.figure(figsize=(15,10))
fig.suptitle('Algorithm Comparison')
ax = fig.add_subplot(111)
plt.boxplot(results)
ax.set_xticklabels(names)
plt.show()
```

Log: 0.844218 (0.013758)
KNN: 0.821769 (0.019432)
CART: 0.774830 (0.028287)
RF: 0.853741 (0.017007)
SVM: 0.838776 (0.003117)
XGB: 0.858503 (0.021252)
[LightGBM] [Info] Number of positive: 213, number of negative: 1110
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000433 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 531
[LightGBM] [Info] Number of data points in the train set: 1323, number of used features: 23
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.160998 -> initscore=-1.650823
[LightGBM] [Info] Start training from score -1.650823
[LightGBM] [Info] Number of positive: 213, number of negative: 1110
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000546 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 533
[LightGBM] [Info] Number of data points in the train set: 1323, number of used features: 23
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.160998 -> initscore=-1.650823
[LightGBM] [Info] Start training from score -1.650823
[LightGBM] [Info] Number of positive: 213, number of negative: 1110
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000219 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 532
[LightGBM] [Info] Number of data points in the train set: 1323, number of used features: 23
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.160998 -> initscore=-1.650823
[LightGBM] [Info] Start training from score -1.650823
[LightGBM] [Info] Number of positive: 213, number of negative: 1110
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000264 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 532
[LightGBM] [Info] Number of data points in the train set: 1323, number of used features: 23
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.160998 -> initscore=-1.650823
[LightGBM] [Info] Start training from score -1.650823
[LightGBM] [Info] Number of positive: 213, number of negative: 1110
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000259 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 531
[LightGBM] [Info] Number of data points in the train set: 1323, number of used features: 23
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.160998 -> initscore=-1.650823
[LightGBM] [Info] Start training from score -1.650823
[LightGBM] [Info] Number of positive: 213, number of negative: 1110
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000241 seconds.
You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

```
[LightGBM] [Info] Total Bins 532
[LightGBM] [Info] Number of data points in the train set: 1323, number of used features: 23
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.160998 -> initscore=-1.650823
[LightGBM] [Info] Start training from score -1.650823
[LightGBM] [Info] Number of positive: 213, number of negative: 1110
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000243
seconds.
```

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

```
[LightGBM] [Info] Total Bins 533
[LightGBM] [Info] Number of data points in the train set: 1323, number of used features: 23
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.160998 -> initscore=-1.650823
[LightGBM] [Info] Start training from score -1.650823
[LightGBM] [Info] Number of positive: 214, number of negative: 1109
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000244
seconds.
```

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

```
[LightGBM] [Info] Total Bins 531
[LightGBM] [Info] Number of data points in the train set: 1323, number of used features: 23
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.161754 -> initscore=-1.645238
[LightGBM] [Info] Start training from score -1.645238
[LightGBM] [Info] Number of positive: 214, number of negative: 1109
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000223
seconds.
```

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

```
[LightGBM] [Info] Total Bins 533
[LightGBM] [Info] Number of data points in the train set: 1323, number of used features: 23
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.161754 -> initscore=-1.645238
[LightGBM] [Info] Start training from score -1.645238
[LightGBM] [Info] Number of positive: 214, number of negative: 1109
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000262
seconds.
```

You can set `force_row_wise=true` to remove the overhead.

And if memory is not enough, you can set `force_col_wise=true`.

```
[LightGBM] [Info] Total Bins 532
[LightGBM] [Info] Number of data points in the train set: 1323, number of used features: 23
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.161754 -> initscore=-1.645238
[LightGBM] [Info] Start training from score -1.645238
LightGBM: 0.863265 (0.013758)
```

Algorithm Comparison

